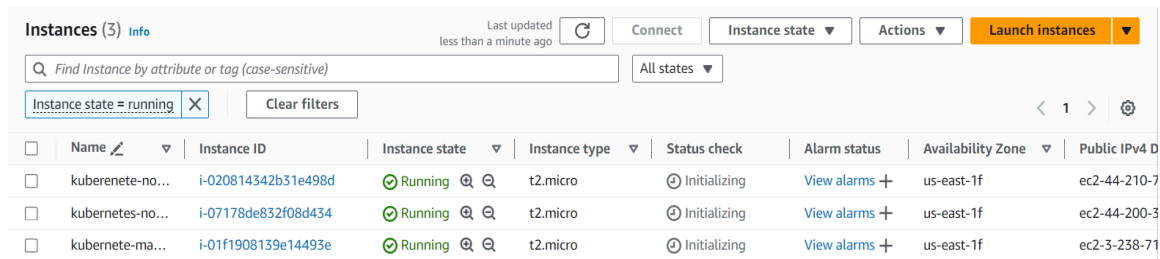


EXP No.:3

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud

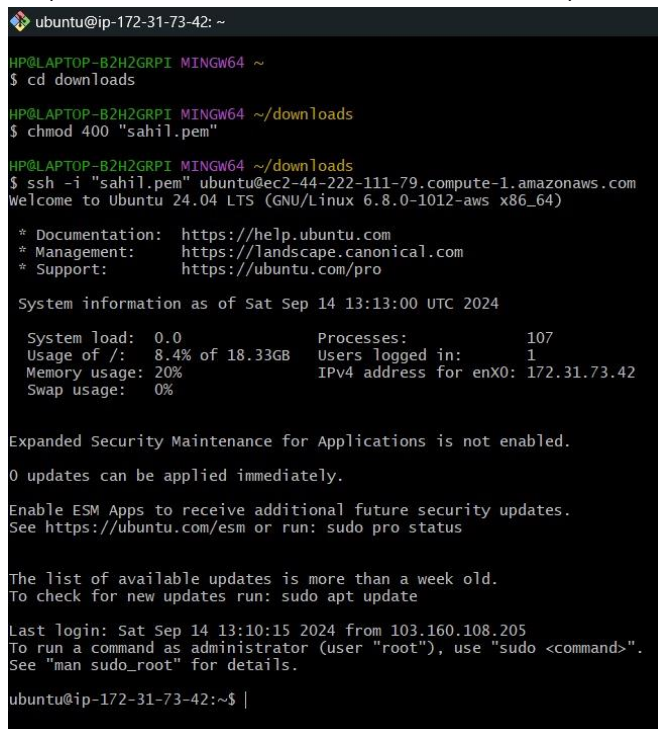
Step 1: Create 3 EC-2 instances with all running on Amazon Linux as OS
Kubernet-master, kubernet-node1, kubernet-node2.



The screenshot shows the AWS Management Console 'Instances' page. It displays three EC2 instances in a 'Running' state. The instances are named 'kubernet-no...', 'kubernet-no...', and 'kubernet-ma...'. They are all using the 't2.micro' instance type and are located in the 'us-east-1f' availability zone. The status check for all instances is 'Initializing'.

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 D
<input type="checkbox"/>	kubernet-no...	i-020814342b31e498d	Running	t2.micro	Initializing	View alarms +	us-east-1f	ec2-44-210-7
<input type="checkbox"/>	kubernet-no...	i-07178de832f08d434	Running	t2.micro	Initializing	View alarms +	us-east-1f	ec2-44-200-3
<input type="checkbox"/>	kubernet-ma...	i-01f1908139e14493e	Running	t2.micro	Initializing	View alarms +	us-east-1f	ec2-3-238-71

Step 2: SSH into all 3 machines each in separate terminal for each instance,



```
ubuntu@ip-172-31-73-42: ~  
HP@LAPTOP-B2H2GRPI MINGW64 ~  
$ cd downloads  
HP@LAPTOP-B2H2GRPI MINGW64 ~/downloads  
$ chmod 400 "sahil.pem"  
HP@LAPTOP-B2H2GRPI MINGW64 ~/downloads  
$ ssh -i "sahil.pem" ubuntu@ec2-44-222-111-79.compute-1.amazonaws.com  
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)  
  
* Documentation:  https://help.ubuntu.com  
* Management:    https://landscape.canonical.com  
* Support:        https://ubuntu.com/pro  
  
System information as of Sat Sep 14 13:13:00 UTC 2024  
  
System load:  0.0          Processes:      107  
Usage of /:   8.4% of 18.3GB Users logged in: 1  
Memory usage: 20%         IPv4 address for enx0: 172.31.73.42  
Swap usage:   0%  
  
Expanded Security Maintenance for Applications is not enabled.  
0 updates can be applied immediately.  
Enable ESM Apps to receive additional future security updates.  
See https://ubuntu.com/esm or run: sudo pro status  
  
The list of available updates is more than a week old.  
To check for new updates run: sudo apt update  
  
Last login: Sat Sep 14 13:10:15 2024 from 103.160.108.205  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
ubuntu@ip-172-31-73-42:~$ |
```

Step 3:Now install docker in all 3 instances;

sudo apt-get install -y docker.io

```
ubuntu@ip-172-31-73-42:~$ sudo apt-get install -y docker.io
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  bridge-utils containerd dns-root-data dnsmasq-base pigz runc ubuntu-fan
Suggested packages:
  ifupdown aufs-tools cgroupfs-mount | cgroup-lite debootstrap docker-buildx docker-compose-v2 docker-doc rinse zfs-fuse | zfsutils
The following NEW packages will be installed:
  bridge-utils containerd dns-root-data dnsmasq-base docker.io pigz runc ubuntu-fan
0 upgraded, 8 newly installed, 0 to remove and 7 not upgraded.
Need to get 76.8 MB of archives.
After this operation, 289 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 pigz amd64 2.8-1 [65.6 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 bridge-utils amd64 1.7.1-1ubuntu2 [33.9 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 runc amd64 1.1.12-0ubuntu3.1 [8599 kB]
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 containerd amd64 1.7.12-0ubuntu4.1 [38.6 MB]
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 dns-root-data all 2023112702~willsync1 [4450 B]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 dnsmasq-base amd64 2.90-2build2 [375 kB]
Get:7 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/universe amd64 docker.io amd64 24.0.7-0ubuntu4.1 [29.1 MB]
Get:8 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 ubuntu-fan all 0.12.16 [35.2 kB]
Fetched 76.8 MB in 1s (74.6 MB/s)
```

Then, Configure File daemon.json;

```
ubuntu@ip-172-31-73-42:~$ cd /etc/docker
ubuntu@ip-172-31-73-42:/etc/docker$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
ubuntu@ip-172-31-73-42:/etc/docker$ |
```

- sudo systemctl enable docker
- sudo systemctl daemon-reload
- sudo systemctl restart docker
- docker -v

```
ubuntu@ip-172-31-73-42:/etc/docker$ sudo systemctl enable docker
ubuntu@ip-172-31-73-42:/etc/docker$ sudo systemctl daemon-reload
ubuntu@ip-172-31-73-42:/etc/docker$ sudo systemctl restart docker
ubuntu@ip-172-31-73-42:/etc/docker$ docker -v
Docker version 24.0.7, build 24.0.7-0ubuntu4.1
ubuntu@ip-172-31-73-42:/etc/docker$ |
```

Step 4: Install Kubernetes in all three instances:

```
[ec2-user@ip-172-31-81-63 docker]$ sudo setenforce 0
[ec2-user@ip-172-31-81-63 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

Add kubernetes repository (paste in terminal);

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes] name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/ enabled=1 gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/r epomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni EOF
```

- `sudo yum update`
- `sudo yum install -y kubelet kubeadm kubectl`
`--disableexcludes=kubernetes`

```
[ec2-user@ip-172-31-81-63 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:01:34 ago on Wed Sep 11 15:39:05 2024.
Dependencies resolved.
=====
Package                                Architecture                               Version
=====
Installing:
kubeadm                                x86_64                                     1.30.4-150500.1.1
kubectl                                x86_64                                     1.30.4-150500.1.1
kubelet                                x86_64                                     1.30.4-150500.1.1
Installing dependencies:
conntrack-tools                        x86_64                                     1.4.6-2.amzn2023.0.2
cri-tools                              x86_64                                     1.30.1-150500.1.1
kubernetes-cni                         x86_64                                     1.4.0-150500.1.1
libnetfilter_cthelper                  x86_64                                     1.0.0-21.amzn2023.0.2
libnetfilter_cttimeout                 x86_64                                     1.0.0-19.amzn2023.0.2
libnetfilter_queue                     x86_64                                     1.0.5-2.amzn2023.0.2
socat                                  x86_64                                     1.7.4.2-1.amzn2023.0.2
Transaction Summary
=====
Install 10 Packages
```

After installing Kubernetes, we need to configure internet options to allow bridging.

- `sudo swapoff -a`
- `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`
- `sudo sysctl -p`

Step 5: Perform this ONLY on the Master machine Initialize kubernetes by

typing below command

- `sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all`

```
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
  --discovery-token-ca-cert-hash sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6cd2d2f6f8c5
[ec2-user@ip-172-31-81-63 docker]$ mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-81-63 docker]$
```

Copy the mkdir and chown commands from the top and execute them

```
mkdir -p $HOME/.kube sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config sudo  
chown $(id -u):$(id -g) $HOME/.kube/config
```

Copy this join link and save it in clipboard (copy from your output as it is different for each instance)

```
kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \
```

```
--discovery-token-ca-cert-hash
```

```
sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6c d2d2f6f8c5
```

Then, add a common networking plugin called flannel file as mentioned in the code.

```
kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
[ec2-user@ip-172-31-81-63 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml  
namespace/kube-flannel created  
clusterrole.rbac.authorization.k8s.io/flannel created  
clusterrolebinding.rbac.authorization.k8s.io/flannel created  
serviceaccount/flannel created  
configmap/kube-flannel-cfg created  
daemonset.apps/kube-flannel-ds created
```

Check the created pod using this command

- kubectl get pods

Step:6. For nodes only;

Use the below command on all 2 node machines

- Sudo yum install iproute-tc-y
- sudo systemctl enable kubelet
- sudo systemctl restart kubelet
- kubeadm join 172.31.81.63:6443 --token zh5jbb.a6ty3eujzc51d15d \

Name: Sahil Motiramani

Div:D15C

Roll No:35

```
--discovery-token-ca-cert-hash  
sha256:0822f656bf52a17a2b6686c123f811306f41495ca650a0aed9bf6cd2d2f6f8 c5
```

Master control nodes;

```
Every 2.0s: kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-81-63.ec2.internal	Ready	control-plane	29m	v1.30.4
ip-172-31-87-137.ec2.internal	Ready	<none>	5m58s	v1.30.4
ip-172-31-92-18.ec2.internal	Ready	<none>	5m53s	v1.30.4